

# Prefetching in RocksDB/Apache Flink for streaming workloads

MASTER THESIS STUDENT

FREDRIK BJÖRKMAN

fbjorkma@kth.se

EXAMINER

PARIS CARBONE

parisc@kth.se

SUPERVISOR

SONIA-FLORENTINA HORCHIDAN

sfhor@kth.se

March 16, 2022

## 1 Preliminary thesis title

Prefetching in RocksDB/Apache Flink for streaming workloads

## 2 Keywords

Apache Flink, RocksDB, Datastreams, Prefetching

## 3 Background

This project will be carried out at KTH in collaboration with Boston University. Modern stream processors rely on embedded key-value stores to manage state that accumulates over long-running computations and exceeds the available memory size. One of these key-value stores is RocksDB [1] which is used by many systems, such as Apache Flink [2]. An important technique in reducing the I/O wait time for disc reads is the use of prefetching data into main memory before it is used. The challenge then arise to determine what data to prefetch, since the memory accesses might not have any obvious patterns. However some unique characteristics arises around streaming state access workloads. The state access for streaming workloads tends to have a high spacial locality [3].

In an article from 2020, the performance of two different state backends, RocksDB and FASTER [4], was evaluated [5]. The report showed that depending on the workload, the state backends would perform differently. When benchmarking with a *COUNT* operation, that represents the entire window as a single value, FASTER, with its support for in place updates, was performing better than RocksDB. On the other hand, when benchmarking with the *RANK* operation, that fetches an entire window's content on trigger, then RocksDB was the one that excelled over FASTER. To quote the authors, this was to showcase that "not one store fits all streaming workloads". They also implemented their own **workload-aware** state backend, that showed a significant boost in performance compared to both FASTER and RocksDB. This to show that being able to adapt based on the workload, would have a great impact on performance.

This project will instead of focusing on adapting the behaviour of the application based on the workload, focus on trying to increase performance by utilizing a prefetching mechanism in the Apache Flink native code. The goal of this project is to leverage this result by designing and implementing a novel prefetching mechanism in RocksDB that is tailored for streaming workloads, that proactively populates the cache with data that has a high probability of being accessed in the near future. To do so, we will need to continuously identify and extract key neighbourhoods from the state access traces of streaming computations.

## 4 Research question and Method

### 4.1 Research question

By designing and implementing an effective prefetch mechanism in Apache Flink when accessing RocksDB. Will the boost in performance of state access on the streaming data be enough to motivate an implementation of this in the official Apache Flink framework?

### 4.2 Hypothesis

The expected outcome of this project is that, if an effective prefetch mechanism is in place, then the improvement in performance would make it worth implementing this mechanism in the official Apache Flink source code.

### 4.3 Research method

This project will have a quantitative approach, where tests will be performed to measure the performance for state access on the streaming data and compare the performance with and without a prefetch mechanism tailored for streaming workloads.

The program used for benchmarking will be Gadget [6]. The benchmarking program is based on a previous report from Boston University. Gadget is a benchmarking program that tests the performance on latency and throughput on a key-value store and in this case RocksDB. Gadget will generate a realistic state access workload that the performance will be measured on.

### 4.4 Objectives and Tasks

The research question can be broken down into the following objectives:

- Implementing a working prefetching mechanism when accessing the stream data states from RocksDB.
- Benchmark the application with and without the prefetching algorithm.

Then following tasks needs to be completed in order to achieve the objectives:

- Implement a simple test program to monitor the communication with RocksDB
- Find a suitable place in the code to implement the prefetching algorithm.
- Try to sort the keys to minimise random access and see how the latency is affected.
- Try to buffer these keys and see how the latency is affected.
- Maintain statistics on the most accessed keys for the workload so we can do smart prefetching.
- Test the prefetching algorithm and evaluate performance.

### 4.5 Ethics and Sustainability

This project aims to improve performance when accessing data from RocksDB and have very little to do with ethics and sustainability. One could argue that if an implementation would be successful and the amount of look-ups in RocksDB would decrease due to data being successfully prefetched, then the system could potentially draw marginally less power.

## 4.6 Limitations

This project will only look at RocksDB as the backend database for Apache Flink. Apache Flink will also be the only stream-processing framework that this project will research. Only streaming data will be researched. There could be a use for a trained machine learning model to decide what data to prefetch, but at this stage in the project it is still unclear if that will be implemented or not.

## 4.7 Risks

As this project will be a quantitative research project and what is to be evaluated is code written by the master thesis student, there is always a risk in implementation errors that can lead to inaccurate conclusions to be drawn. This will hopefully be avoided by a careful well thought through implementation and potential code reviews.

# 5 Evaluation and News Value

## 5.1 Evaluation

This project can be evaluated by creating a benchmark program that uses the same operations on the same data stream and comparing the current implementation with the implementation that contains the prefetch mechanism.

## 5.2 Expected scientific results

Currently there is no support for a prefetching mechanism in between Apache Flink and RocksDB. By performing this project (and creating a proof on concept), it may lead to decide if it is worth to implement this in the product to increase performance. The proof of concept and the research behind it might also map out the way this could be implemented in the product.

## 5.3 The work's innovation/news value

The ones interested in this work will most likely be the ones that are developing and using Apache Flink. If the developers would implement this feature it would hopefully lead to a boost in performance for the application.

This might also be of interest for those who look for similar solutions between a streaming framework and a state backend database.

# 6 Pre-study

This project, as said in Section 3, will be built upon previous work done by Vasiliki Kalavri and John Liagouris, *"In support of workload-aware streaming state management"* [5]. The pre-study will originate from that work to get an understanding of the background to this project. After that there will be a lot of work to read documentation and source code for both Apache Flink and RocksDB. On top of that, information need to be acquired regarding how the implantation of the prefetching algorithm is going to take place. The literature study will need to be conducted in parallel with the implementation throughout a large part of this project as more knowledge about Apache Flink and RocksDB are acquired. Documentation and source code will be read about Gadget [6] to get an understanding on how the implementations will be benchmarked and how to configure Gadget so that it benchmarks what is supposed to be benchmarked.

## 7 Conditions and Schedule

### 7.1 Resources

The project will be able to use an already build testbed built by Boston University for the performance testing on state access for streaming workloads. There are also available data sets for streaming data that can be used for testing of the performance. Both KTH and Boston University have expressed that they are able to assist with computational resources. (Update) See 4.3 for some specification on what is going to be used from Boston university.

### 7.2 Milestone charts

This is the estimated timeline for tasks in project:

Task(s)	Dates
Experimenting with Flink/RockDB	7/2 - 18/2
Dig into source code	
Write individual plan	14/2 - 15/2
Implement test program	
Find suitable place in the native Flink code for the prefetching to happen	21/2 - 4/3
Write background	7/3 - 18/3
Update Individual plan	15/3 - 16/3
Implement batching of keys in native Flink to minimize the number of accesses to RocksDB.	
Learn about gadget and benchmark batching vs current Flink	21/3 - 25/3
Implement sorting of keys to minimise random access.	
Take notes for the method chapter.	28/3 - 14/4
Easter (Will not work on thesis)	15/4 - 18/4
Study for exam (Will not work on thesis)	18/4 - 22/4
Implement a way to maintain statistics for accessed keys to do smart prefetching.	
Take notes for the method chapter.	25/4 - 6/5
Run benchmarks	9/5 - 20/5
Work on Method and Results	23/5 - 3/6
Work on Discussion, Conclusion, Future work, Abstracts and Introduction	6/6 - 10/6
Prepare for Final presentation	13/6 - 17/6
Finalize report	20/6 - 1/7

This is the estimated timeline for deadlines in project:

Deadline	Date
Individual plan	15/2
Updated individual plan	16/3
Draft of Background	18/3
Implementation done	6/5
Tests done	20/5
Draft of Method and Results	3/6
Draft of report as a whole	10/6
Final presentation	17/6
Final report done	1/7

## References

- [1] D. B. et al, “RocksDB,” *GitHub*, 12 2021, [Online; accessed 13. Dec. 2021]. [Online]. Available: <https://github.com/facebook/rocksdb/wiki/RocksDB-Overview>
- [2] “Apache Flink: What is Apache Flink? — Applications,” 9 2021, [Online; accessed 14. Dec. 2021]. [Online]. Available: <https://flink.apache.org/flink-applications.html>
- [3] S. Somogyi, T. Wenisch, A. Ailamaki, B. Falsafi, and A. Moshovos, “Spatial memory streaming,” in *33rd International Symposium on Computer Architecture (ISCA’06)*, vol. 34, no. 2. IEEE Computer Society, 2006. doi: 10.1109/ISCA.2006.38 pp. 252–263.
- [4] B. Chandramouli, G. Prasaad, D. Kossmann, J. Levandoski, J. Hunter, and M. Barnett, “Faster: A concurrent key-value store with in-place updates,” in *Proceedings of the 2018 International Conference on Management of Data*, 2018, pp. 275–290.
- [5] V. Kalavri and J. Liagouris, “In support of workload-aware streaming state management,” in *12th USENIX Workshop on Hot Topics in Storage and File Systems (HotStorage 20)*. USENIX Association, 7 2020. [Online]. Available: <https://www.usenix.org/conference/hotstorage20/presentation/kalavri>
- [6] Casp-Systems-Bu, “Gadget,” 3 2022, [Online; accessed 15. Mar. 2022]. [Online]. Available: <https://github.com/CASP-Systems-BU/Gadget>

## Acronyms

**KTH** KTH Royal Institute of Technology